

Welcome back! If you found this question useful,  
don't forget to vote both the question and the answers up.

[close this message](#)

## calculating nucleotide frequency per column

Asked 4 years ago   Active 1 year, 4 months ago   Viewed 2k times



I have some sequences shown below

5



1



CAGGTAGCC  
CCGGTCAGA  
AGGGTTTGA  
TTGGTGAGG  
CAAGTATGA  
ACTGTATGC  
CTGGTAACC  
TATGTA CTG  
GCTGTGAGA  
CAGGTGGGC  
TCAGTGAGA  
GGGGTGAGT  
TGGGTATGT  
GAGGTGAGA  
CAGGTGGAG

Each line has 9 nucleotides. Consider it to be 9 columns. I want to calculate the nucleotide frequency of each nucleotide for each of the 9 columns. For example 1st column will have these bases C,C,A,T,C,A etc Out put should be something like this

A	0.25	0.34	0.56	0.43	0.00	0.90	0.45	0.34	0.31
C	0.45	0.40	0.90	0.00	0.40	0.90	0.30	0.25	0.2
G	0.00	0.00	0.34	1.00	0.30	0.30	0.35	0.90	0.1
T	0.24	0.56	0.00	0.00	1.00	0.34	0.45	0.35	0.36

**Note**, I just made up the numbers to show you the output file format

Welcome back! If you found this question useful,  
don't forget to vote both the question and the answers up.

close this message

Share Edit Follow Flag

edited Jul 22 '20 at 15:55

asked Dec 14 '17 at 16:40



M\_

5,713 2 17 38



user3138373

400 1 3 12

▲ I think a package in Bioconductor can do this, I don't recall exactly one. Has you searched there? Or what have you tried to calculate this? The original file is a plain text or a fasta file ( I don't recall if it is only for fasta or for other file types)? – [llrs](#) Dec 14 '17 at 16:57

▲ Original file is a text file only. I am thinking of using awk(associative arrays) to calculate it, still writing the script though – [user3138373](#) Dec 14 '17 at 17:02

## 7 Answers

Active

Oldest

Votes



For anyone interested in doing this from any sort of alignment file, I've implemented a [position frequency matrix function in AlignBuddy](#).

3

input:



```
15 9
seq_1  CAGGTAGCC
seq_2  CCGGTCAGA
seq_3  AGGGTTTGA
seq_4  TTGGTGAGG
seq_5  CAAGTATGA
seq_6  ACTGTATGC
seq_7  CTGGTAACC
seq_8  TATGTACTG
seq_9  GCTGTGAGA
seq_10 CAGGTGGGC
seq_11 TCAGTGAGA
seq_12 GGGGTGAGT
seq_13 TGGGTATGT
seq_14 GAGGTGAGA
seq_15 CAGGTGGAG
```

Welcome back! If you found this question useful,  
don't forget to vote both the question and the answers up.

[close this message](#)

```
:$ alignbuddy input.phy --pos_freq_mat
```

Output:

```
### Alignment 1 ###
A      0.133  0.400  0.133  0.000  0.000  0.400  0.467  0.067  0.400
C      0.400  0.267  0.000  0.000  0.000  0.067  0.067  0.133  0.267
G      0.200  0.200  0.667  1.000  0.000  0.467  0.200  0.733  0.200
T      0.267  0.133  0.200  0.000  1.000  0.067  0.267  0.067  0.133
```

[Share](#) [Edit](#) [Follow](#) [Flag](#)

[edited Dec 21 '17 at 15:15](#)

[answered Dec 16 '17 at 17:08](#)



**Steve Bond**

**176** 1 6

awk

5

```
awk '{L=length($1);for(i=1;i<=L;i++) {B=substr($1,i,1);T[i][B]++;}} END{for(BI=0;BI<4;BI++) {B=(BI==0?"A":(BI==1?"C":
(BI==2?"G":"T")));printf("%s",B); for(i in T) {tot=0.0;for(B2 in T[i]){tot+=T[i][B2];}printf("\t%.2f", (T[i][B]/tot));}
printf("\n");}}' input.txt
```



```
A  0.13  0.40  0.13  0.00  0.00  0.40  0.47  0.07  0.40
C  0.40  0.27  0.00  0.00  0.00  0.07  0.07  0.13  0.27
G  0.20  0.20  0.67  1.00  0.00  0.47  0.20  0.73  0.20
T  0.27  0.13  0.20  0.00  1.00  0.07  0.27  0.07  0.13
```

Or, expanded for clarity:

```
awk '{
    L=length($1);
    for(i=1;i<=L;i++) {
        B=substr($1,i,1);
```

Welcome back! If you found this question useful,  
don't forget to vote both the question and the answers up.

[close this message](#)

```
B=(BI==0?"A":(BI==1?"C":(BI==2?"G":"T")));
printf("%s",B);
for(i in T) {
    tot=0.0;
    for(B2 in T[i]){
        tot+=T[i][B2];
    }
    printf("\t%.2f", (T[i][B]/tot));
}
printf("\n");
}
```

} 'input.txt'

Share Edit Follow Flag

edited Dec 16 '17 at 16:34

answered Dec 15 '17 at 11:21



terdon

7,802 3 14 40



Pierre

1,448 5 11



Can we do same for taking dinucleotide at time (eg; AA, TT, CG etc); the output matrix would reduce to half? – [kashiff007](#) Aug 4 at 7:32



4



#Make a list of hashes

```
hl = []
```

```
for i in range(9):
```

```
    hl.append({'A': 0, 'C': 0, 'G': 0, 'T': 0})
```

```
f = open("foo.txt") # CHANGE ME
```

```
nLines = 0
```

```
for line in f:
```

```
    for idx, c in enumerate(line.strip()):
```

```
        hl[idx][c] += 1
```

```
        nLines += 1
```

Welcome back! If you found this question useful,  
don't forget to vote both the question and the answers up.

[close this message](#)

```
print("{}\t{}".format(char, "\t".join(["{:0.2f}".format(x[char]/nLines) for
x in hl])))
```

The output of your example is then:

A	0.13	0.40	0.13	0.00	0.00	0.40	0.47	0.07	0.40
C	0.40	0.27	0.00	0.00	0.00	0.07	0.07	0.13	0.27
G	0.20	0.20	0.67	1.00	0.00	0.47	0.20	0.73	0.20
T	0.27	0.13	0.20	0.00	1.00	0.07	0.27	0.07	0.13

[Share](#) [Edit](#) [Follow](#) [Flag](#)

[edited Dec 15 '17 at 10:22](#)

[answered Dec 14 '17 at 21:01](#)



[terdon](#)

**7,802** 3 14 40



[Devon Ryan](#) ♦

**18.9k** 2 23 50



how to change enumerate if I want to take dinucleotides at time (eg; AA, TT, CG etc); the output matrix would reduce to half? – [kashiff007](#) Aug 4 at 7:37

Here's a Perl approach:

3

```
#!/usr/bin/env perl
use strict;
my %counts;

## Read the input file line by line
while (my $line = <>) {
    print;
    ## remove trailing '\n' characters
    chomp $line;
    ## split the line into an array at every character
    my @columns=split(/./, $line);
    ## iterate over the array from the first to the last position
    for my $i (0..$#columns){
```

Welcome back! If you found this question useful,  
don't forget to vote both the question and the answers up.

close this message

```

    $counts{$nt}[$i]++;
  }
}
## Iterate over the counts hash
for my $nt (sort keys(%counts)){
  print "$nt\t";
  ## dereference the array stored in the hash
  my @countsForThisNt = @{$counts{$nt}};
  ## iterate over the counts for each position for this nt
  for (my $l=0;$l<=$#countsForThisNt;$l++) {#
    ## If the value for this position isn't defined,
    ## set it to 0.
    $countsForThisNt[$l]=0;
    ## Print all the things
    printf "%.2f\t", $countsForThisNt[$l]/$. , $l;
  }
  print "\n";
}

```

Save the script somewhere in your PATH, make it executable and run:

```

$ foo.pl file
A  0.13    0.40    0.13    0.00    0.00    0.40    0.47    0.07    0.40
C  0.40    0.27    0.00    0.00    0.00    0.07    0.07    0.13    0.27
G  0.20    0.20    0.67    1.00    0.00    0.47    0.20    0.73    0.20
T  0.27    0.13    0.20    0.00    1.00    0.07    0.27    0.07    0.13

```

Alternatively, [if you're into the whole brevity thing](#), and enjoy some golfing, here's the same thing as a one-liner:

```

perl -ne 'chomp;@F=split(/);for$i(0..$#F){$k{$F[$i]}[$i]++}{for $nt(sort keys(%k)){print"$nt\t";for$i(0..$#{k{$nt}}){$g=k{$nt}[$i]||0; printf "%.2f\t",$g/$.}print"\n";}}' file

```

Share Edit Follow Flag

edited Dec 15 '17 at 10:12

answered Dec 15 '17 at 10:01

 [terdon](#)

Welcome back! If you found this question useful,  
don't forget to vote both the question and the answers up.

close this message

2

which make this manual work unnecessary. Here's a more idiomatic solution:

```
sequences = readLines('sequences.txt')
bases_matrix = do.call(rbind, strsplit(sequences, ''))

apply(bases_matrix, 2L, function (col) {
  str = DNASTring(paste(col, collapse = ''))
  letterFrequency(str, letters = 'ACGT', OR = 0L, as.prob = TRUE)
})
```

This uses the same Bioconductor packages. Since this is such a simple problem, it can also be written without Bioconductor:

```
bases = strsplit(sequences, '')
# Use a data.frame here so we can use factors in the next step:
# R does not support matrices of factors. Ugh.
bases_by_column = setNames(do.call(rbind.data.frame, bases),
  seq_along(bases[[1L]]))
# Ensure that every column will be a complete set of ACGT frequencies
bases_by_column = lapply(bases_by_column, factor, c('A', 'C', 'G', 'T'))
sapply(lapply(bases_by_columns, table), prop.table)
```

Using modern R idioms from the 'magrittr' package, I'd write this as a pipeline; this very directly shows the sequence of transformations.

```
do.call(rbind.data.frame, bases) %>%
  setNames(seq_along(bases[[1L]])) %>%
  lapply(factor, c('A', 'C', 'G', 'T')) %>%
  lapply(table) %>%
  sapply(prop.table)
```

Share Edit Follow Flag

edited Jul 22 '20 at 13:58

answered Dec 21 '17 at 15:08










Konrad Rudolph

4,665 10 38

Welcome back! If you found this question useful,  
don't forget to vote both the question and the answers up.

close this message

- 1  @b.nota Nothing wrong with "it works". But this being a Q&A site, I think we should definitely also teach best practices. In fact, there's a problem in bioinformatics with this "it works" attitude, which leads to a proliferation of a lot of truly atrocious code. "But it works". Yeah, just about ... as long as you don't nudge it. – [Konrad Rudolph](#) Dec 21 '17 at 16:04 
- 1  I don't completely agree, no offense. With many bioinformaticians the BIO part is much more important than the best script practices. At the company where I work, they don't care about these best script practises, but do want someone that understands e.g., what a tissue resident memory cell is. That's just my point of view. – [benn](#) Dec 21 '17 at 16:09 
- 1  @b.nota It's akin to lab (or generally science) best practices, to be honest: Many people use the same argument to justify sloppy lab work or dodgy use of statistics (p-hacking etc). At the moment bad software engineering is still somewhat accepted but I'm confident that this acceptance will lessen as the field matures. Just to clarify: this doesn't so much concern your code, which works fine. But there's a lot of genuinely bad, unmaintainable code in circulation in bioinformatics *tools* that costs money and time to fix because of bad practices. – [Konrad Rudolph](#) Dec 21 '17 at 16:13 
- 1  Yeah I agree with that part. I am PhD in biology, and learned to code with a java course here and a course in R there. I work with R almost daily, but unfortunately I still try to avoid `apply` functions (I do try them more often, but still the old school java double for loop is more intuitive for me). – [benn](#) Dec 21 '17 at 16:19

Here an example in R using `Biostrings` and `letterFrequency` (as suggested by [Devon Ryan](#)).

2

```
library(Biostrings)
```

```
data <- read.table("DNA.txt", stringsAsFactors = F)
```

```
new <- matrix(nrow = 9, ncol = 15)
```

```
for(i in 1:9){
  for(j in 1:15){
    new[i,j] <- substring(data[j,], i, i)
  }
}
```

```
countTable <- matrix(nrow = 9, ncol = 4)
for(i in 1:9){
  columnSeq <- DNASTringSet(paste0(new[i,], collapse = ""))
  columnCounts <- letterFrequency(columnSeq, letters = "ACGT", OR = 0)
  countTable[i,] <- columnCounts
}
```



Welcome back! If you found this question useful,  
don't forget to vote both the question and the answers up.

[close this message](#)

```

      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9]
A 0.13 0.40 0.13  0  0 0.40 0.47 0.07 0.40
C 0.40 0.27 0.00  0  0 0.07 0.07 0.13 0.27
G 0.20 0.20 0.67  1  0 0.47 0.20 0.73 0.20
T 0.27 0.13 0.20  0  1 0.07 0.27 0.07 0.13

```

Share Edit Follow Flag

[edited Dec 15 '17 at 19:59](#)

[answered Dec 15 '17 at 9:07](#)



[benn](#)

**3,511** 6 24

▲ The whole new matrix generation can be written in a single line as `do.call(rbind, strsplit(data, ''))`. The subsequent for loop can likewise be replaced. – [Konrad Rudolph](#) Dec 21 '17 at 15:10 ✎

▲ Yes, there are many roads that lead to Rome. – [benn](#) Dec 21 '17 at 15:55

▲ Here is how you get the distributions by column using the `TraMineR` R package.

1 `library(TraMineR)`

▼ `sts.data <- c(
 "CAGGTAGCC",
 "CCGGTCAGA",
 "AGGGTTTGA",
 "TTGGTGAGG",
 "CAAGTATGA",
 "ACTGTATGC",
 "CTGGTAACC",
 "TATGTACTG",
 "GCTGTGAGA",
 "CAGGTGGGC",
 "TCAGTGAGA",
 "GGGGTGAGT",
 "TGGGTATGT"`

Welcome back! If you found this question useful,  
don't forget to vote both the question and the answers up.

[close this message](#)

```
seqstatd(seq)
```

and here is the outcome of the `seqstatd` function

```
[State frequencies]
  [1] [2] [3] [4] [5] [6] [7] [8] [9]
A 0.13 0.40 0.13  0  0 0.400 0.467 0.067 0.40
C 0.40 0.27 0.00  0  0 0.067 0.067 0.133 0.27
G 0.20 0.20 0.67  1  0 0.467 0.200 0.733 0.20
T 0.27 0.13 0.20  0  1 0.067 0.267 0.067 0.13
```

```
[Valid states]
  [1] [2] [3] [4] [5] [6] [7] [8] [9]
N 15 15 15 15 15 15 15 15 15
```

```
[Entropy index]
  [1] [2] [3] [4] [5] [6] [7] [8] [9]
H 0.94 0.94 0.62  0  0 0.78 0.87 0.62 0.94
```

[Share](#) [Edit](#) [Follow](#) [Flag](#)

answered Dec 21 '17 at 10:51



[Gilbert](#)

**111** 2